
JUDICIOUS BATCHING: THE IMPACT OF BATCH ORDERING IMAGING TESTS IN THE EMERGENCY DEPARTMENT

A PREPRINT

Jacob Jameson

Interfaculty Initiative in Health Policy
Harvard University
Cambridge, MA
jacobjameson@g.harvard.edu

Soroush Saghafian

Kennedy School of Government
Harvard University
Cambridge, MA
soroush_saghafian@hks.harvard.edu

Robert Huckman

Harvard Business School
Harvard University
Boston, MA
rhuckman@hbs.edu

Nicole Hodgson

Department of Emergency Medicine
Mayo Clinic
Phoenix, AZ
Hodgson.Nicole@mayo.edu

Abstract

We use practice variation across physicians to uncover the role of test-ordering on care delivery in the ED. Using records of over 45,000 Mayo Clinic emergency department visits...

Keywords Emergency Department · Operational Efficiency · Diagnostic Testing

1 Introduction

Healthcare delivery, particularly in the emergency department (ED), requires a delicate balance between ensuring optimal patient outcomes and optimizing resource utilization. Achieving these twin goals necessitates timely and accurate diagnosis, which enables prompt and appropriate treatment, improving patient prognosis and reducing the likelihood of adverse events. Efficient patient discharge from the ED is crucial for alleviating overcrowding, a severe issue associated with higher complication rates and increased mortality (bernstein2009?). In this complex environment, the availability and performance of diagnostic tests, especially imaging studies, play a pivotal role in shaping both patient care and operational efficiency (naseim2015?).

A critical question in ED management pertains to whether physicians should batch order diagnostic imaging tests or order them sequentially. This decision represents a fundamental tradeoff between potentially reducing patient length of stay and risking over-testing. Batch ordering, where multiple imaging tests are ordered simultaneously, may expedite the diagnostic process but can lead to unnecessary tests, increased costs, patient anxiety, and potential harm from follow-up of false-positive results (koch2018?). Conversely, a more sequential approach might reduce unnecessary testing but could extend a patient's time in the ED, potentially contributing to overcrowding. Striking the right balance between comprehensive evaluation and efficient resource use is crucial, yet research on optimal test ordering strategies in the ED remains limited (saghafian2015?).

In this paper, we leverage data from over 45,000 ED patient visits to quantify the benefits and consequences of batching versus sequentially ordering diagnostic imaging tests. Our focus on imaging tests is particularly relevant given their significant impact on patient flow, resource utilization, and potential health risks from radiation exposure. We employ an empirical strategy that exploits the random assignment of patients to ED physicians who differ in their propensity to batch-order imaging tests. This quasi-experimental design allows

us to overcome potential selection bias and move closer to identifying the causal impact of batch ordering on key outcomes such as patient length of stay, resource utilization, and the likelihood of return visits.

To capture physician tendencies towards batch ordering, we introduce a novel measure using a leave-out, residualized approach based on each physician’s historical ordering patterns for all other patients seen during the study period. This tendency measure strongly predicts actual batch ordering behavior while remaining uncorrelated with patient and ED visit characteristics, providing a robust instrument for our analysis. Our findings reveal significant practice variation among physicians and demonstrate substantial impacts on ED operations and patient care trajectories.

We begin by evaluating the reduced-form effects of ED provider batch-ordering tendency on downstream patient outcomes and ED turnaround time. Our results indicate that practice variation, as captured by physician batch-ordering tendency, has large and significant consequences. For instance, being treated by a provider in the top decile of the batching tendency distribution, compared to one in the bottom decile, results in additional diagnostic imaging tests per 100 patient encounters, even after controlling for the physician’s underlying propensity to test. This variation in practice patterns raises important questions about the optimal use of diagnostic resources in emergency care settings.

To address potential confounding factors, we employ a placebo exercise to demonstrate that the observed effects are primarily due to batching behavior rather than other aspects of physician care that might correlate with batching tendency. We find no significant effects of physician batch tendency on key outcomes for a “placebo sample” of patients visiting the ED for conditions that rarely result in batched orders. This strengthens our confidence in the specificity of the batching effect on patient outcomes and resource utilization.

Our research design closely approximates a randomized controlled trial that assigns patients to batch-ordering or sequential-ordering arms, leveraging the unique features of the ED setting. In this environment, patients have no discretion over provider selection, and in our specific ED, the random assignment of patients to physicians further mitigates selection issues. The wide variation in batch-ordering practices among physicians within the same hospital, despite following identical guidelines, provides the necessary variation for our analysis. Moreover, the typically short, one-off nature of ED patient-physician interactions constrains decision-making to a more limited, observable set of choices compared to other healthcare settings.

By exploiting this practice variation in ED settings, we aim to isolate the impact of test batching from other potential channels that might influence length of stay and patient outcomes in different healthcare contexts. This approach allows us to more accurately identify the causal impact of batch-ordering diagnostic imaging tests on patient outcomes and resource utilization. Importantly, our study examines batching decisions made within the bounds of normal clinical judgment and practice norms, rather than focusing on deviations from clinical guidelines or instances of substandard care.

The implications of this research extend beyond academic interest, offering potential insights for ED management, healthcare policy, and the design of clinical decision support systems. Understanding the true impacts of batch ordering practices could inform more efficient resource allocation strategies, improve patient flow, and ultimately enhance the quality of emergency care delivery. The remainder of this paper is structured as follows: The next section describes our data source and outlines the baseline sample. Section II describes the theoretical and conceptual model. Section IV details our empirical strategy and its accompanying identifying assumptions. Section V presents our results, including both reduced-form effects and instrumental variable analyses. Section V also discusses the implications of our findings for batch-ordering policies and ED management. The final section concludes with a summary of our contributions and directions for future research.

2 Theoretical Model

Emergency departments (EDs) are complex environments where critical decisions about patient care must be made rapidly and under conditions of uncertainty. Our theoretical model aims to capture the nuanced dynamics of diagnostic decision-making in this high-stakes setting, with a particular focus on the practice of batch ordering imaging tests.

In a typical ED scenario, patients arrive with varying levels of acuity and a wide range of presenting symptoms. Upon arrival, each patient i is assigned to a physician j , often based on availability rather than specific matching criteria. The physician’s primary task is to diagnose and treat the patient efficiently and effectively, a process that frequently involves ordering diagnostic tests, including imaging studies.

Central to our model is the concept of the patient’s underlying condition, denoted by τ_i . This represents the true state of the patient’s health, which is initially unknown to the treating physician. The physician’s goal is to uncover τ_i as accurately and quickly as possible to provide appropriate care. However, the path to this discovery is not straightforward and involves critical decisions about diagnostic testing.

The physician’s approach to ordering tests is influenced by various factors, including their individual tendencies and experiences. We represent this as Z_j , the physician’s propensity to batch order tests. This tendency can be thought of as a stable characteristic of the physician, shaped by their training, past experiences, and personal practice style. The decision to batch order tests for a given patient, represented by X_{ij} , is a function of both the physician’s tendency Z_j and the specifics of the patient’s presentation. Batch ordering involves requesting multiple imaging tests simultaneously, rather than in a sequential manner. This decision has cascading effects on the patient’s journey through the ED.

Figure 1: Conceptual Model




Figure 1 illustrates the causal relationships between physician tendency, batch ordering decision, test volume, and patient outcomes.

As illustrated in Figure 1, when a physician chooses to batch order ($X_{ij} = 1$), it directly influences the total volume of imaging tests performed, M_{ij} . Our model posits two key mechanisms for this relationship:

1. **Precautionary ordering:** In the face of diagnostic uncertainty, physicians may opt for a comprehensive initial workup, potentially including tests that might not have been ordered in a sequential approach.
2. **Reduced information gain:** By ordering tests simultaneously, physicians forego the opportunity to use the results of initial tests to inform the necessity of subsequent ones. This can lead to a higher overall test volume compared to a sequential approach where each test order is informed by preceding results.

The total volume of tests, M_{ij} , in turn, impacts key patient outcomes, Y_{ij} , particularly length of stay (LOS) and the probability of a 72-hour return visit with admission. The relationship between test volume and these outcomes is complex. On one hand, more tests generally require more time, potentially increasing LOS. However, if the physician’s initial diagnostic uncertainty is high, ordering more tests upfront might expedite the diagnostic process and lead to quicker treatment decisions.

Importantly, our model also allows for a direct effect of the batch ordering decision on outcomes, independent of its effect through test volume. This direct effect could represent, for example, the efficiency gains from parallel processing of multiple tests, which might reduce LOS even as the total number of tests increases. The model acknowledges the presence of unobserved confounding factors, U_{ij} , which might influence both the decision to batch order and patient outcomes. These could include factors like the physician’s assessment of the patient’s condition severity or subtle clinical cues not captured in standard data.

To formalize these relationships, we express them mathematically as follows:

$$p(X_{ij}|Z_j) = f(Z_j, U_{ij}) \quad (1)$$

$$M_{ij} = g(X_{ij}, U_{ij}) \quad (2)$$

$$Y_{ij} = h(X_{ij}, M_{ij}, U_{ij}) \quad (3)$$

To address potential endogeneity, our empirical strategy leverages the physician’s general tendency to batch order, Z_j , as an instrument for the actual batching decision, X_{ij} . This approach allows us to estimate the causal effects of batch ordering on test volume and patient outcomes, even in the presence of unobserved confounders.

Our empirical analysis focuses on estimating these key relationships:

1. The effect of physician tendency on batching decisions: $\frac{\partial X_{ij}}{\partial Z_j}$
2. The effect of batching on test volume: $\frac{\partial M_{ij}}{\partial X_{ij}}$
3. The total effect of batching on patient outcomes: $\frac{\partial Y_{ij}}{\partial X_{ij}}$

In summary, our theoretical model provides a framework for understanding the complex decision-making process in ED diagnostic testing. It highlights the potential trade-offs involved in batch ordering – between comprehensive initial assessment and efficient resource use, between rapid testing and informed sequential decision-making. By formalizing these relationships, we set the stage for our empirical analysis, which aims to quantify these effects and provide insights for optimizing ED operations and patient care.

3 Data and Definitions

Our study was conducted in the Emergency Department (ED) of the Mayo Clinic of Arizona, a tertiary care hospital without obstetrical services, an inpatient pediatrics unit, or a trauma designation. During the study period, the ED recorded approximately 43,000 visits per year, managed across 26 treatment rooms and up to 9 hallway spaces. The department is exclusively staffed by board-eligible or board-certified emergency physicians (EPs), with rotating residents overseeing about 10% of patient volume. Physicians operate in a unique workflow that includes staggered 8.5-hour shifts and a rotational patient assignment system, minimizing potential selection bias in patient encounters.

We conducted a retrospective review of comprehensive ED operational data from 10/6/2018 through 12/31/2019, coinciding with the initiation of a new electronic medical record. The dataset includes detailed patient demographics, chief complaints, vital signs, emergency severity index (ESI), length of stay (LOS), and resource utilization metrics. This period was chosen to provide a robust data set while excluding the influence of the coronavirus pandemic.

3.1 Sample Construction and Summary Statistics

Our research design focuses on adult visits to the Mayo Clinic of Arizona ED. From an initial dataset of approximately 48,000 visits during the study period, we applied several exclusion criteria to improve statistical power and ensure the validity of our physician batch tendency instrument. We excluded encounters with rare chief complaints (fewer than 1000 total encounters) and complaints where batch ordering occurs in less than 5% of cases. Examples of excluded complaints include Upper Respiratory Symptoms and Urinary Complaints. This exclusion strategy is unlikely to introduce selection bias as long as physician test batching tendency is not strongly correlated with diagnosing behavior across broad complaint categories. To estimate a precise measure of physician-level batch tendency, we further restricted our sample to encounters involving full-time physicians who treat over 500 ED cases per year. These criteria resulted in a final analytical sample of 16,361 patient encounters.

Table 1 provides a comprehensive overview of ED characteristics, patient demographics, and medical tests in our study site. The ED manages an average volume of 24.3 patients, indicating a significant yet manageable

workload. The patient population exhibits a range of critical conditions, with 18.8% presenting as tachycardic, 9.51% as tachypneic, 2.43% with fever, and 1.57% as hypotensive. Diagnostic procedures in our sample heavily rely on laboratory tests, used in 78.5% of encounters. Imaging plays a crucial role, with X-rays ordered in 48.1% of visits, non-contrast CT scans in 23.1%, and contrast CT scans in 20.5%. Ultrasounds are used in 12.5% of encounters. Notably, 6.32% of visits involve batch ordering of imaging tests, a key focus of our study.

Table 1: Summary Statistics of Emergency Department Encounters

	Mean	Q1	Median	Q3
Emergency Department Characteristics				
Patients in ED	24.3	16	25	33
Tachycardic	0.188			
Tachypneic	0.095			
Febrile	0.024			
Hypotensive	0.016			
ESI	2.74	2	3	3
Complaint: Abdominal Complaints	0.169			
Complaint: Extremity Complaints	0.143			
Complaint: Chest Pain	0.096			
Complaint: Neurological Issue	0.095			
Complaint: Gastrointestinal Issues	0.090			
Patient Characteristics				
Arrival Age	58.3	44	61	75
Race: White	0.886			
Race: Black	0.041			
Race: Asian	0.030			
Gender: Female	0.545			
Tests				
X-Ray	0.481			
Ultrasound	0.125			
Non-Contrast CT	0.231			
Contrast CT	0.205			
Lab	0.785			
Imaging Tests were Batch Ordered	0.063			

This table reports summary statistics for the baseline sample of emergency department visits during the study period described in the text. Vital signs were categorized as follows: tachycardia (pulse more significant than 100), tachypnea (respiratory rate greater than 20), fever (temperature greater than 38°C), and hypotension (systolic blood pressure less than 90).

This sample composition and these descriptive statistics provide essential context for our subsequent analyses of batch ordering practices and their impacts on ED efficiency and patient outcomes. The diversity in patient characteristics, presenting complaints, and diagnostic approaches underscores the complexity of ED operations and the potential significance of test ordering strategies in this setting.

3.2 Variable Definitions

Our explanatory variable in the IV analysis, $Batched_i$, is an indicator for whether patient i has their tests batch-ordered at their ED encounter. While the patient could decide not to undergo the tests ordered by the physician, this is rare in practice. Below, we detail our primary outcomes: (a) ED length of stay (LOS), and (b) total imaging volume, and (c) 72-hour return with admission.

(a) *ED length of stay (LOS)*– ED-LOS is a critical measure of efficiency and patient throughput in emergency care settings. It is defined as the duration from a patient’s arrival to the ED until their departure, whether by discharge or admission to the hospital. The hypothesis is that batch testing may lead to shorter ED-LOS, potentially improving patient throughput and reducing crowding.

(b) *Resource Utilization*– Resource utilization in the ED typically refers to the extent of medical services and interventions a patient receives. In this study, we quantify resource utilization by the total number of distinct diagnostic tests ordered per patient during their ED stay. This encompasses both initial and subsequent tests. The hypothesis is that batch testing may lead to variations in the number of tests ordered, potentially influencing the overall healthcare expenditure and efficiency.

(c) *72 Hour Return with Admission*– This is a binary variable indicating whether the patient returns to the ED within 72 hours of their initial visit and is admitted to the hospital. This outcome is used to assess the quality of care provided during the initial ED visit. The hypothesis is that batch testing may lead to variations in the quality of care provided, potentially influencing the likelihood of a return visit.

3.2.1 Batching

We define “batching” in line with standard emergency medicine practices. Batching occurs when a physician simultaneously orders a comprehensive set of diagnostic tests, typically covering a broad range of potential diagnoses. This contrasts with sequential ordering, where tests are ordered sequentially based on the information obtained from each test as needed.

We operationalize batching as occurring when multiple diagnostic imaging tests are ordered within a 5-minute window. In Section 4.4, sensitivity analyses on this cutoff point showed that our results are robust to this definition. Each imaging test (e.g., X-ray, CT scan, Ultrasound) is considered a separate, distinct test for our study. Therefore, a batch in our study consists of two or more distinct imaging tests.

4 Empirical Strategy

Our empirical strategy closely follows the literature that relies on quasi-random assignment of agents to cases, often referred to as the “judges design.” Papers in this literature typically exploit variation in the sentencing leniency of judges who work in the same court. Similarly, we explore batching variation across physicians who work in the same emergency department. In its reduced form, under the assumption of quasi-random assignment, this approach allows researchers to identify the causal effect of being assigned to different types of physicians. Under additional assumptions, an instrumental variable approach identifies the causal effect of a given medical decision. We employ both approaches and lay out their details in the next subsections.

4.1 Institutional Details on Patient-Physician Assignment

Contrary to most healthcare settings where patients exhibit choice, they are predominantly passive in their physician assignment in the ED. In most EDs, however, physicians have discretion in picking their patients. In contrast, patients arriving at the Mayo Clinic ED are randomly assigned to physicians via a rotational patient assignment algorithm ([traub2016emergency?](#)), which removes potential selection bias concerns for our analyses. In essence, barring arrival time and shift-level variation, the physician-to-patient matching can be deemed random. Table 2 displays that patient encounters (regarding chief complaints and emergency severity) are equitably distributed across physicians within our study’s cohort.

4.2 Batch Tendency Construction

To measure physician batch tendency, we use the physician’s residualized leave-out average batch rate. This measure is derived from two steps following the approaches taken by ([doyle2015measuring?](#)), ([dobbie2018effects?](#)), and ([eichmeyer2022pathways?](#)). First, we obtain residuals from a regression model, which includes all ED encounters in our sample period.

$$Batched_{i,t} = \alpha_0 + \alpha_{ym} + \alpha_{dt} + \alpha_{complaint_esi} + \alpha_{lab} + \varepsilon_{i,t} \quad (4)$$

Where $Batched_{i,t}$ is a dummy variable equal to one if patient i had their imaging tests batch ordered on encounter that took place on date t . Fixed effects include year-month fixed effects, α_{ym} , to control for time and seasonal variation in batching, such as hospital-specific policies (e.g. initiatives to eliminate excess

Table 2: Balancing Test: Wald Test for Equality of Means

Chief Complaints	Frequency	F-Statistic	$Pr(> F)$
Abdominal Complaints	6232	2.587	0.108
Back or Flank Pain	2552	1.637	0.201
Chest Pain	3525	0.407	0.524
Extremity Complaints	5265	1.847	0.174
Falls, Motor Vehicle Crashes, Assaults, and Trauma	2381	0.023	0.880
Gastrointestinal Issues	3323	0.105	0.746
Neurological Issue	3495	0.135	0.713
Shortness of Breath	2966	1.324	0.250
Skin Complaints	2178	0.383	0.536
Upper Respiratory Symptoms	1917	0.017	0.896
Emergency Severity	Frequency	F-Statistic	$Pr(> F)$
ESI 1 or 2	13914	0.011	0.915
ESI 3, 4, or 5	29386	0.010	0.921

Table 2 reports the results of a Wald test which was conducted to assess the balance of chief complaints across providers in our dataset. A balanced distribution implies that complaints and severity are evenly distributed across providers, which we expect to be the case due to randomization. The Wald F-statistic and p-value are reported. Robust standard errors (type HC1) were used to account for potential heteroscedasticity in the data.

testing) or seasonality in ED visits. We also control for “shift-level” variations that include both physician scheduling and patient arrival with day of week-time of day fixed effects, α_{dt} . Chief complaint by severity fixed effects, $\alpha_{complaint}$, were also included to increase precision. Finally, a binary variable for whether or not laboratory tests were ordered, α_{lab} , was included to account for the complexity of the case. As stated earlier, these controls are more than what is required for our quasi-random assignment assumption. Under the assumption that we have captured the observables under which quasi-random assignment occurs in the ED, the unexplained variation— the physician’s contribution— resides in the error term, $\varepsilon_{i,t}$.

In step two, the tendency measure for patient i seen by physician j is computed as the average residual across all other patients seen by the physician that year:

$$Tendency_{i,j}^{phys} = \frac{1}{N_{-i,j}} \sum_{i' \in \{J \setminus i\}} \hat{\varepsilon}_{i'} \quad (5)$$

where $\hat{\varepsilon}_{i'} = Batch_{i'} - \hat{Batch}_{i'}$ is the residual from equation (1); J is the set of all ED encounters treated by physician j ; and $N_{-i,j} = |\{J \setminus i\}|$, the number of cases that physician has seen that year, excluding patient i . This leave-out mean eliminates the mechanical bias that stems from patient i ’s own case entering into the instrument. The measure is interpreted as the average (leave-out) batch rate of patient i ’s physician, relative to other physicians in that hospital-year-month, hospital-day of week-time of day.

We document that the Mayo Clinic ED physicians exhibit wide, systematic variation in their propensity to batch order diagnostic tests. Figure 2 graphs the frequency with which physicians batch order across different chief complaints, highlighting that the variation in batching differs starkly. We note that there is high correlation between

Table illustrates the “first stage” results, emphasizing the robust relevance of the instrument across various controls. Being assigned to a 10pp higher batch-tendency physician is associated with a 17.6pp increase in the probability of having tests batch-ordered in the ED. The F-statistic is 21.6 when all controls and fixed effects are included. The coefficient is greater than one because all emergency visits are used to construct the tendency instrument, while the first stage is calculated using the baseline sample only.

To estimate the reduced-form effects of being treated by a batch-preferring physician, we estimate the following equation:

$$Y_i = \mu_0 + \mu_1 Tendency_{i,j}^{phys} + \gamma X_i + \nu_i \quad (6)$$

Figure 2: Variation in Physician Batching




Figure 2 illuminates the marked differences among physicians in their propensity to batch-order diagnostic tests.

Table 3: Comparison of First Stage Estimates

	Coefficient		
	(1)	(2)	(3)
Batch Tendency	1.92*** (0.07)	1.91*** (0.06)	1.76*** (1.76)
Day of Week-Time of Day FE	✓	✓	✓
Month of Year FE	✓	✓	✓
Complaint/Severity FE		✓	✓
Laboratory Tests Ordered			✓
F Statistic	9.55	17.86	21.59
N	16,361	16,361	16,361

Estimates of the first stage for the baseline sample described in the text. Seasonality shift fixed effects include Year-Month and Hospital-Day of week-Hour of day fixed effects. Chief complaint comes from the cleaned complaint that the patient came in with at the initial encounter. Column 3 corresponds to the baseline controls. Robust standard errors are clustered at the physician level.

*** $p < 0.001$, ** $p < 0.01$, * $p < 0.05$.

This reduced form allows us to check the strength of our instrument. To study the effects of test batching in the ED, we employ a multi-stage approach using our baseline sample:

$$Tests_i = \alpha_0 + \alpha_1 Batched_i + \delta X_i + \eta_i \quad (7)$$

$$Y_i = \beta_0 + \beta_1 Batched_i + \beta_2 Tests_i + \theta X_i + \varepsilon_i \quad (8)$$

Where $Tests_i$ is the number of imaging tests ordered, and Y_i represents our patient outcomes of interest: length of stay and 72-hour return with admission. X_i is the same set of control variables as in the reduced-form approach. The $Batched_i$ variable suffers from potential endogeneity concerns. For example, unobserved injury severity may correlate with both the need for multiple tests and patient outcomes. To address this, we instrument $Batched_i$ with the assigned physician j 's underlying tendency to batch, $Tendency_{i,j}^{phys}$. We use a

causal mediation analysis to decompose the total effect of batching into its direct effect on outcomes and its indirect effect through the number of tests ordered. This approach employs bootstrapping to estimate confidence intervals for the direct and indirect effects. We cluster standard errors at the physician level to account for potential correlation in outcomes among patients treated by the same physician, and to properly adjust for the shared physician-level instrument across patients.

4.3 Identifying Assumptions

The reduced-form approach delivers an unbiased estimate of the causal effect of being treated by a higher tendency to batch physician, since assignment of patients to ED physicians is random, conditional on seasonality and shift (“conditional independence”). The residualization in equation (4) controls for more controls than required to achieve quasi-random assignment; they are included for statistical precision in measuring physician tendency to batch.

Our instrumental variable approach, which aims to recover the causal effect of having diagnostic tests batch ordered, relies on three additional assumptions: relevance, exclusion, and monotonicity. We reported a strong first stage (i.e., relevance) at the end of the previous Section. The exclusion restriction requires that the instrument must influence the outcome of interest only through its effect on test batching. This is perhaps our strongest assumption and is at its core, untestable. However, several features of the ED setting suggest that such violation may likely only have a small impact and may be less concerning than in other health care settings. First, unlike in primary care settings, where the patient and primary care provider have many repeat encounters, the scope of what the emergency physician can do to impact medium-term outcomes is limited and well-observed by the researcher. Second, any violation of the exclusion restriction needs to directly affect the specific outcome of interest. The channel by which ED physicians can influence length of stay relative outcomes is likely through testing and diagnosis. Nevertheless, we take this assumption seriously and perform a placebo check in Section 5.2 as well as various robustness checks in Section 6.

Finally, the monotonicity assumption is necessary for interpreting the coefficient estimates obtained from the IV approach as Local Average Treatment Effects (LATEs) if there are heterogeneous treatment effects. It requires that any patient who is (not) batched by a sequencer (batcher) would also (not) be batched by a batcher (sequencer) physician. The literature leveraging the judges design typically performs two informal tests for its implications. The first one provides that the first stage should be weakly positive for all subsamples ((dobbie2018effects?)). The second implication asserts that the instrument constructed by leaving out a particular subsample has predictive power over that same left-out subsample ((bhuller2020incarceration?)). Appendix Table ?? presents both of these tests in the two columns for various subsamples of interest.

5 Results

5.1 Reduced-Form Results

In this section, we explore the causal influence of physician batch tendency on total imaging volume in the emergency department. We posit that while batch tendency directly influences the practice of batch ordering tests, both batch tendency and batch ordering are concurrently influenced by a physician’s testing inclination. Given that testing inclination directly affects primary outcomes, we include it as a control variable in our regression models to mitigate its confounding effects.

To quantify a physician’s testing inclination, we employ a similar approach to that used in measuring physician batch tendency. Specifically, we calculate the physician’s residualized leave-out average imaging test order rate, which serves as a proxy for their propensity to order tests. It is important to note the strong positive correlation between batch tendency and testing inclination, suggesting that physicians with a higher propensity to test may also exhibit a higher tendency to batch orders, potentially as a consequence of their testing strategies. Neglecting to account for testing inclination could lead to overestimated effects of batch tendency due to omitted variable bias.

Reduced-form regression results in Table ?? reveal the effect the average total effect of batch tendency on the number of distinct imaging tests ordered for our final sample. We hypothesize that while batch ordering may streamline the testing process, resulting in a quicker completion of a given number of tests, it simultaneously appears to lead to an increase in the total number of tests ordered. The increased testing volume, in turn, is associated with an extended LOS and possibly lower rate of return with admission.

Table 4: Reduced Form Results: Effect of Batch Tendency on the Number of Unique Imaging Tests Ordered

	Coefficient		
	(1)	(2)	(3)
Batch Tendency	4.543*** (0.532)	1.11*** (0.124)	1.11* (0.447)
<i>Baseline Controls</i>	✓	✓	✓
<i>Testing Inclination</i>		✓	✓
<i>Any Imaging Ordered</i>			✓
R^2	0.203	0.207	0.599
N	16,361	16,361	16,361

Estimates of the reduced form for the baseline sample and baseline controls described in the text. Seasonality shift fixed effects include Year-Month and Hospital-Day of week-Hour of day fixed effects. Chief complaint comes from the cleaned complaint that the patient came in with at the initial encounter. Robust standard errors are clustered at the physician level.

*** $p < 0.001$, ** $p < 0.01$, * $p < 0.05$.

To determine whether or not the effect of batch tendency on LOS is mediated by the number of tests ordered, we estimate a mediation analysis using the approach outlined by (imai2010general?). Our analysis, scrutinizes the pathways through which physician batch tendency influences patient length of stay (LOS) in the emergency department. The analysis delineates both direct effects, which capture the influence of batch tendency on LOS without intermediaries, and indirect effects, which operate through the mediator of test ordering volume. Our findings underscore a significant Average Direct Effect (ADE) of batch tendency on LOS, estimated at -0.15 ($p = 0.046$).

5.2 Placebo Check

In this section we investigate whether the reduced-form effects observed in Section 5.1 are due to differences in batch rates across providers or due to other provider differences correlated with batch tendency. We start by studying reduced-form effects among patients with complaints that are never batched, as a “placebo/falsification check.” By way of example, consider a patient who arrives at the ED with a urinary tract infection—a condition for which patients rarely undergo imaging testing. For such patients, we should expect to see no impact of batch tendency only if high batching and low batching physicians do not systematically differ in other dimensions of care relevant to patient outcomes. Conversely, if we do find a reduced-form effect for these patients, then high batch tendency physicians must systematically differ from low batch tendency physicians in other dimensions of care, beyond batching.

To that end, we restrict attention to ED visits for complaints where batching occurs no more than 10 percent of the time (recall that our baseline sample only includes complaints with a >10 percent batching rate). We estimate a reduced-form regression of each main outcome on physician batch tendency for the subsample, following equation (3). The results of this exercise are displayed in Appendix Table 1. They show that in contrast to results for our main sample, the association between physician tendency to batch and a given outcome is statistically indistinguishable from zero and much smaller in magnitude for the samples of patients who visit the ED with health conditions that are rarely batched.

5.3 Instrumental Variables Results

This section presents our causal analysis of batch ordering imaging tests in the ED, focusing on its effects on resource utilization and patient outcomes. Our analysis reveals nuanced impacts of batch ordering, with significant implications for ED management and patient care.

5.3.1 Effect of Batch Ordering on Imaging Tests and Patient Outcomes

Table 5 presents the results of our two-stage least squares (2SLS) analysis and subsequent causal mediation analysis for our analytical sample.

Our 2SLS estimates reveal a significant positive effect of batch ordering on the number of imaging tests performed. Specifically, batch ordering leads to an increase of 0.699 imaging tests (95% CI: [0.138, 1.260], p

Table 5: Combined 2SLS and Causal Mediation Analysis Results

	Number of Imaging Tests	Log LOS	72-hr Return with Admission
<i>2SLS Results</i>			
Batched	0.699** (0.286)	-0.539 (1.534)	-0.036 (0.034)
<i>Baseline Controls</i>	✓	✓	✓
<i>Testing Inclination</i>	✓	✓	✓
<i>Imaging Ordered</i>	✓	✓	✓
<i>Causal Mediation Analysis</i>			
ACME		0.083* [0.017, 0.15]	-0.003** [-0.005, 0.00]
ADE		-0.565 [-3.608, 2.56]	-0.034 [-0.101, 0.03]
Total Effect		-0.481 [-3.549, 2.60]	-0.037 [-0.104, 0.03]

Note: ** $p < 0.05$, * $p < 0.1$. Full controls included. Estimates from 2SLS and causal mediation analysis, presenting impacts on number of tests, length of stay, and 72-hour return rates.

< 0.05) per patient encounter. This finding confirms our hypothesis that batch ordering results in a higher volume of diagnostic imaging compared to a sequential approach. The magnitude of this effect suggests that for every 100 batch orders, approximately 70 additional imaging tests are performed that would not have occurred under a sequential ordering strategy.

This excess in imaging tests can be attributed to the fundamental difference between batch and sequential ordering approaches. In a sequential approach, physicians can incorporate the information gained from initial tests to inform decisions about subsequent testing, potentially reducing unnecessary examinations. Batch ordering, by contrast, commits to multiple tests upfront, without the benefit of this stepwise information gain.

5.3.2 Mediation Analysis: Imaging Volume as a Mediator

Our causal mediation analysis provides crucial insights into how the increased imaging volume affects key patient outcomes. For length of stay (LOS), we find a significant Average Causal Mediation Effect (ACME) of 0.0834 (95% CI: [0.0171, 0.15], $p < 0.1$). This positive ACME indicates that batch ordering indirectly extends ED stays through increased testing. Notably, while the Average Direct Effect (ADE) of batch ordering on LOS is negative (-0.565), it is not statistically significant. These findings challenge the conventional wisdom that batch ordering is a time-saving measure. While there may be some efficiency gains in the testing process itself (as suggested by the negative, albeit non-significant, ADE), these are more than offset by the additional time required to perform and interpret the excess tests generated by batch ordering. The net result is a tendency towards longer patient stays in the ED.

Interestingly, our analysis of 72-hour return rates with admission reveals a contrasting effect. We observe a small but statistically significant negative ACME of -0.00274 (95% CI: [-0.00550, 0.00], $p < 0.05$). This suggests that the increased imaging from batch ordering may slightly reduce the probability of patients returning to the ED within 72 hours and requiring admission. This finding indicates that the “shotgun” approach of comprehensive initial testing may have some benefits in terms of thorough patient evaluation, potentially catching issues that might otherwise lead to return visits. These results paint a nuanced picture of the impacts of batch ordering in the ED. On one hand, it leads to increased resource utilization and longer patient stays, which could strain ED capacity and potentially contribute to overcrowding. On the other hand, it may result in more comprehensive initial evaluations that slightly reduce short-term return visits.

The complexity of these findings underscores the need for careful consideration in ED test ordering strategies. While batch ordering may offer some benefits in terms of comprehensive initial evaluations, it also carries significant costs in terms of resource utilization and patient flow. ED managers and policymakers must weigh these trade-offs carefully when developing guidelines for diagnostic testing. These results also highlight the importance of considering both direct and indirect effects when evaluating ED processes. The apparent efficiency of batch ordering in the testing process itself may be misleading if one does not account for the downstream consequences of increased test volume. Given the varying impacts of batch ordering across

different outcomes, it is likely that the appropriateness of this practice may depend on specific clinical scenarios or patient characteristics. To further explore this possibility and provide more targeted insights for ED management, we next conduct a subgroup analysis examining the effects of batch ordering across different presenting complaints.

5.3.3 Subgroup Analysis by Chief Complaint

Our subgroup analysis reveals important differences in the effects of batch ordering across various presenting complaints, offering valuable insights for optimizing diagnostic strategies in the ED. Table 6 presents these results.

Table 6: Subgroup Analysis Results by Chief Complaint

Chief Complaint	Batched Coefficient	ACME LOS	ADE LOS	ACME 72hr Return	ADE 72hr Return
Extremity Complaints	0.709* (0.411)	0.133* [-0.030, 0.290]	0.268 [-3.640, 3.870]	-0.003 [-0.012, 0.000]	-0.090 [-0.214, 0.040]
Falls/MVCs/Trauma	0.268 (0.422)	0.040 [-0.103, 0.180]	-0.019 [-3.109, 2.960]	-0.002 [-0.009, 0.000]	0.037 [-0.055, 0.130]
Dizziness/Syncope	1.572** (0.701)	0.193** [0.015, 0.390]	-1.353 [-4.043, 1.370]	-0.011** [-0.025, 0.000]	-0.023 [-0.270, 0.210]
Neurological Issue	1.652 (1.983)	0.139 [-0.242, 0.520]	-7.353 [-16.602, 2.500]	0.003 [-0.020, 0.030]	-0.220 [-0.987, 0.560]
Fatigue and Weakness	0.793 (0.686)	0.072 [-0.059, 0.220]	-0.286 [-1.422, 0.940]	-0.002 [-0.013, 0.000]	0.010 [-0.102, 0.120]
Fevers/Sweats/Chills	-0.181 (0.858)	-0.025 [-0.250, 0.200]	-1.086 [-3.079, 1.000]	0.002 [-0.020, 0.030]	0.061 [-0.336, 0.430]

Note: * $p < 0.1$, ** $p < 0.05$. ACME: Average Causal Mediation Effect, ADE: Average Direct Effect. LOS is measured in log-transformed hours. 72hr Return is the probability of return within 72 hours with admission.

For extremity complaints, which range from minor injuries to potentially severe conditions like deep vein thrombosis, batch ordering significantly increases imaging volume (coefficient = 0.709, $p < 0.1$) and indirectly extends ED length of stay (ACME = 0.133, 95% CI: [-0.030, 0.290], $p < 0.1$). However, this approach also shows a trend towards reduced 72-hour returns (ACME = -0.003, 95% CI: [-0.012, 0.000]). These findings suggest that while batch ordering for extremity complaints leads to more resource utilization and longer stays, it may be preventing missed diagnoses of serious conditions like fractures or vascular issues. For extremity complaints, a targeted batch ordering approach may be beneficial. Clinicians should consider batch ordering for patients presenting with symptoms suggestive of more serious conditions (e.g., suspected fractures, vascular compromises) or for those with complex medical histories. For minor injuries or low-risk presentations, a sequential approach might be more appropriate to balance resource use and patient flow.

Dizziness and syncope complaints demonstrate even stronger effects of batch ordering. The significant increase in imaging volume (coefficient = 1.572, $p < 0.05$) translates to longer ED stays (ACME = 0.193, 95% CI: [0.015, 0.390], $p < 0.05$) but also significantly reduces 72-hour returns (ACME = -0.011, 95% CI: [-0.025, 0.000], $p < 0.05$). This pattern likely reflects the complex nature of these complaints, which can stem from benign causes or indicate serious underlying conditions like stroke or cardiac issues. For dizziness and syncope, the reduction in return visits justifies a more liberal use of batch ordering. The potential to identify serious underlying conditions outweighs the increased length of stay. However, clinicians should still use clinical judgment to identify low-risk patients (e.g., young patients with clear peripheral vertigo) who might benefit from a more sequential approach.

For other complaint categories, including falls/trauma, neurological issues, and general symptoms like fatigue or fever, the benefits of batch ordering are less clear-cut. In these cases, the lack of significant effects on return visits suggests that a more selective, sequential approach to test ordering may be more appropriate. For these complaints, we recommend a more individualized approach. Clinicians should rely on thorough initial assessments and risk stratification to determine the need for comprehensive imaging. Decision support tools incorporating factors like patient age, comorbidities, and specific symptoms could guide this process.

It is important to consider several key tradeoffs when implementing batch ordering in the ED:

1. Resource Utilization vs. Patient Safety: While batch ordering increases resource use and ED length of stay, the potential reduction in return visits for certain complaints suggests improved diagnostic

accuracy. EDs must weigh the upfront costs against the long-term benefits of preventing missed diagnoses and return visits.

2. **ED Flow vs. Comprehensive Care:** Longer ED stays due to batch ordering could impact overall ED flow. However, for complaints like dizziness/syncope, this tradeoff appears justified by the reduced return rates. EDs should consider implementing parallel processes (e.g., rapid results reporting, streamlined discharge planning) to mitigate the impact on patient flow.
3. **Cost Considerations:** While our study doesn't directly address costs, the increase in imaging tests from batch ordering likely increases immediate healthcare costs. However, this may be offset by reduced costs from prevented return visits and earlier diagnosis of serious conditions. A comprehensive cost-benefit analysis would be valuable for future research.

5.4 Robustness

5.4.1 Definition of Batching

The 2SLS results presented in the previous sections are robust to several alternative specifications probing the key batching definition. Our primary definition operationalizes batching as the ordering of multiple diagnostic tests within a 5-minute window. This approach is grounded in standard emergency medicine practices and offers a clear distinction between batching and non-batching behaviors. To ensure the robustness of our findings, we performed several sensitivity analyses.

- (1) **Variation in the Time Window for Batching:** Firstly, we varied the time window for what constitutes a batch. The original 5-minute window was extended and contracted to understand its impact on the study's outcomes. Specifically, we considered scenarios where a batch is defined as two or more tests ordered within windows of 1 minute to 10 minutes. This range allowed us to capture a broader spectrum of batching behaviors and test the sensitivity of our results to different batching definitions.
- (2) **Refinement of Batching Definition:** Secondly, we refined our batching definition based on the sequence and timing of test orders. In the initial analysis, if two tests were ordered upfront and another test ordered later, it was counted as a batch. We modified this definition to consider a set of tests as a batch only if all tests ordered during a patient encounter came from that initial batch order. This adjustment ensures that our batching definition more accurately reflects a comprehensive diagnostic effort at the outset of patient care, rather than incremental decision-making.

The key finding from these robustness checks is the consistent impact of batching on key outcomes across all variations. As seen in Appendix Table 2 and 3, altering the time window for batching, whether narrowing it to as little as 1 minute or expanding it to 10 minutes, did not qualitatively change the results. Similarly, refining the definition of batching to consider only those tests ordered in the initial batch also had negligible impact on the study's main findings. These results lend further credibility to our initial findings, demonstrating that our conclusions about the impact of batching on patient outcomes in the ED are not sensitive to the specific operational definition of batching.

5.4.2 External Validity

To assess the generalizability of our findings and ensure their robustness, we conducted an additional analysis using data from the Emergency Department of Massachusetts General Hospital (MGH). This high-volume ED sees approximately 110,000 patient encounters annually, providing a diverse and extensive dataset for comparison. Our analysis utilized 12 months of ED visit data from MGH, encompassing a wide range of patient presentations, diagnoses, and outcomes. This dataset provides a valuable contrast to our primary study site, offering insights into how batch ordering practices and their effects may vary in a different clinical environment.

It is important to note that the MGH ED does not employ the same quasi-random assignment of patients to physicians as our primary study site. This difference in patient allocation necessitated an alternative analytical approach. To address potential selection bias and create a fair comparison between batched and non-batched cases, we employed a propensity score matching technique, which involves:

- **Score Calculation:** We calculated propensity scores for each patient encounter based on key variables including: Chief complaint, Severity (as measured by the Emergency Severity Index or an equivalent triage score), Time of day, Day of week, and Patient demographics (age, gender).

- **Matching:** Using these propensity scores, we matched patients who received batch-ordered imaging tests with those who did not. We employed a nearest-neighbor matching algorithm with a caliper width of 0.2 standard deviations of the logit of the propensity score, which is widely accepted in the literature as a balance between matching quality and number of matched pairs.
- **Balance Checking:** After matching, we conducted balance checks to ensure that the matched groups were comparable across all covariates used in the propensity score calculation. This included visual inspection of covariate distributions and statistical tests (e.g., standardized mean differences) to confirm adequate balance.
- **Outcome Comparison:** With our matched sample, we compared outcomes between the batched and non-batched groups, focusing on the same key metrics as our primary analysis: length of stay, number of imaging tests performed, and 72-hour return rates with admission.

This approach allows us to estimate the effect of batch ordering in a setting where randomization is not feasible, by creating comparable groups of batched and non-batched cases. While this method cannot fully replicate the strengths of our primary quasi-experimental design, it provides a valuable check on the external validity of our findings. By comparing the results from this analysis to our primary findings, we can assess whether the effects of batch ordering on ED outcomes are consistent across different hospital settings and patient populations. This comparison enhances the generalizability of our conclusions and provides additional context for our recommendations on ED imaging practices. The results of this external validation analysis are presented in Appendix Table XX.

6 Conclusion